

# K Nearest

*Eoin Flynn*

*18 March 2018*

Bond University  
Data Science

# Contents

<b>Functions</b>	<b>3</b>
<b>Data</b>	<b>6</b>
<b>Model</b>	<b>6</b>
Model Creation . . . . .	7
Model Comparison . . . . .	11
Model Comparison Results . . . . .	13

## Functions

This section will hold all of the functions that will be used throughout this markdown.

```
# Change rows to factors
setRowAsFactor <- function(dataset, columns) {
  for (column in columns) {
    dataset[, column] <- as.factor(dataset[, column])
  }
  return(dataset)
}

# Get the models predictions
getPredictions <- function(model, dataset, outcomeColumn, type = "prob") {
  suppressMessages(library(ROCR))

  probability <- as.numeric(predict(model, newdata = customerDataset,
    type = "prob")[, 2])
  predictions <- prediction(probability, outcomeColumn)

  return(predictions)
}

# Get the model AUC and return it
getModelAUC <- function(model, dataset, outcomeColumn, type = "prob") {
  suppressMessages(library(ROCR))

  prediction <- getPredictions(model = model, dataset = dataset,
    outcomeColumn = outcomeColumn)
  auc <- performance(prediction, "auc")@y.values

  return(auc)
}

# Plot ROC curves
plotROCCurves <- function(model1Prediction, model2Prediction,
  main, model1Colour = "#009900", model2Colour = "#FF8000",
  model1Name, model2Name, legendLocation = "bottomright") {
  suppressMessages(library(ROCR))

  model1Performance <- performance(model1Prediction, "tpr",
    "fpr")
  model2Performance <- performance(model2Prediction, "tpr",
    "fpr")

  plot(model1Performance, main = main, col = model1Colour,
    print.auc = TRUE)
  plot(model2Performance, add = T, col = model2Colour)
```

```

    legend(legendLocation, legend = paste(rep(c(model1Name, model2Name))),
           col = c(model1Colour, model2Colour), cex = 0.8, fill = c(model1Colour,
           model2Colour))
}

# Calculate AUC. Returns as a decimal
getAUC <- function(outcomeColumn, dataset, model, oneClass, zeroClass) {
  suppressMessages(library(ModelMetrics))
  prediction_df <- createPrediction_df(model, dataset, oneClass = oneClass,
    zeroClass = zeroClass)
  auc <- auc(outcomeColumn, prediction_df$classification)

  return(auc)
}

# Create a confusion matrix. Returns a confusion matrix
createConfusionMatrix <- function(model, dataset, outcomeColumn,
  oneClass, zeroClass) {
  suppressMessages(library(caret))
  prediction_df <- createPrediction_df(model, dataset, oneClass = oneClass,
    zeroClass = zeroClass)
  userConfusionMatrix <- table(prediction_df$classification,
    outcomeColumn)

  return(userConfusionMatrix)
}

# Create a new customer for prediction. Returns a dataframe
createCustomer <- function(originalDataset, gender, SeniorCitizen,
  Partner, Dependents, tenure, PhoneService, MultipleLines,
  InternetService, OnlineSecurity, OnlineBackup, DeviceProtection,
  TechSupport, StreamingTV, StreamingMovies, Contract, PaperlessBilling,
  PaymentMethod, MontlyCharges, TotalCharges, Churn) {
  # Create a copy of the original dataset and keep one row that
  # will be overridden with the new data.
  newCustomer <- customerDataset[1, ]

  newCustomer$gender <- gender
  newCustomer$SeniorCitizen <- SeniorCitizen
  newCustomer$Partner <- Partner
  newCustomer$Dependents <- Dependents
  newCustomer$tenure <- tenure
  newCustomer$PhoneService <- PhoneService
  newCustomer$MultipleLines <- MultipleLines
  newCustomer$InternetService <- InternetService
  newCustomer$OnlineSecurity <- OnlineSecurity
  newCustomer$OnlineBackup <- OnlineBackup
  newCustomer$DeviceProtection <- DeviceProtection
  newCustomer$TechSupport <- TechSupport
  newCustomer$StreamingTV <- StreamingTV
  newCustomer$StreamingMovies <- StreamingMovies
  newCustomer$Contract <- Contract
  newCustomer$PaperlessBilling <- PaperlessBilling

```

```

newCustomer$PaymentMethod <- PaymentMethod
newCustomer$MonthlyCharges <- MonthlyCharges
newCustomer$TotalCharges <- TotalCharges
newCustomer$Churn <- Churn

# Convert fields that are factors
newCustomer <- setRowAsFactor(newCustomer, c("gender", "SeniorCitizen",
      "Partner", "Dependents", "PhoneService", "MultipleLines",
      "InternetService", "OnlineSecurity", "OnlineBackup",
      "DeviceProtection", "TechSupport", "StreamingTV", "StreamingMovies",
      "Contract", "PaperlessBilling", "PaymentMethod", "Churn"))

return(newCustomer)
}

# Gets a dataframe from a locally hosted MySQL server.
# Returns a dataframe
loadDataframeFromMySQL <- function(user, password, host = "localhost",
  dbname, statement, port = 3306) {
  suppressMessages(library(RMySQL))

  # Connect to the server
  dataBase <- dbConnect(MySQL(), user = user, password = password,
    host = host, dbname = dbname, port = port)
  # Retrieve the info from the specified server
  dataframe <- dbGetQuery(dataBase, statement)
  # Close the connection to the server
  dbDisconnect(dataBase)

  return(dataframe)
}

# Returns a predictive model
predictiveModel <- function(formula, dataset, method, neighbours = 1:10,
  metric = "Accuracy", trControl) {
  suppressMessages(library(caret))
  model <- train(formula, data = dataset, method = method,
    tuneGrid = expand.grid(.k = neighbours), metric = metric,
    trControl = trControl)

  return(model)
}

# Returns a train control object to be fed into a predictive
# model
trainControlObject <- function(method = "repeatedcv", number = 10,
  repeats) {
  suppressMessages(library(caret))
  object <- trainControl(method = method, number = number,
    repeats = repeats)

  return(object)
}

```

```

}

# Create and return a dataframe of the classifications and
# their probabilities
createPrediction_df <- function(model, dataset, predictionType = "prob",
  oneClass, zeroClass) {
  # Run the prediction
  prediction <- suppressWarnings(predict(model, dataset, type = predictionType))
  # Convert to a dataframe
  prediction_df <- data.frame(prediction)
  # Rename the column to reference easier
  colnames(prediction_df) <- c("NegativeProb", "PositiveProb")
  # Add a row for the classification
  prediction_df$classification <- rep(zeroClass, nrow(prediction_df))
  # Convert all probabilities above 0.5 to be the affirmative
  # class
  prediction_df$classification[prediction_df$PositiveProb >
    0.5] <- oneClass
  prediction_df$classification <- as.factor(prediction_df$classification)

  return(prediction_df)
}

```

## Data

In this section we will load in our data and do some basic data exploration.

```

customerDataset <- loadDataframeFromMySQL(user = "root", password = "A13337995",
  dbname = "world", statement = "Select * from world.customerChurn")

# Loop through and change all relevant rows to factors and
# returns the dataset post modification
customerDataset <- setRowAsFactor(customerDataset, c("gender",
  "SeniorCitizen", "Partner", "Dependents", "PhoneService",
  "MultipleLines", "InternetService", "OnlineSecurity", "OnlineBackup",
  "DeviceProtection", "TechSupport", "StreamingTV", "StreamingMovies",
  "Contract", "PaperlessBilling", "PaymentMethod", "Churn"))

# Drop the columns that will not be needed
customerDataset <- customerDataset[, -which(names(customerDataset) %in%
  c("customerID"))]

```

## Model

We will create two models, one using all variables from the customer dataset, the other using the top three variables outlined in our decision tree paper. The top three variables are the type on contract, the type of internet service, and the customer's tenure. We will test our models between 3 and 45 neighbours as 1 and 2 neighbours typically have little value and more than 45 may make the model prone to overfitting. After the models have been created we will examine their overall accuracy and how accurate they are when they

predict that a customer will churn. Once we have established which model has the greatest accuracy for the business's intended use we will examine if any other number of neighbours close to the one chosen by R would provide a better model for the intended use case.

## Model Creation

We will now create out two models and explore their results.

```
library(caret)

## Loading required package: lattice
## Loading required package: ggplot2
set.seed(12216)

allVariableModel <- predictiveModel(formula = Churn ~ ., dataset = customerDataset,
  method = "knn", neighbours = 3:45, trControl = trainControlObject(repeats = 5))

topThreeModel <- predictiveModel(formula = Churn ~ Contract +
  InternetService + tenure, dataset = customerDataset, method = "knn",
  neighbours = 3:45, trControl = trainControlObject(repeats = 5))
allVariableModel

## k-Nearest Neighbors
##
## 7032 samples
## 19 predictor
## 2 classes: 'No', 'Yes'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 5 times)
## Summary of sample sizes: 6329, 6329, 6329, 6329, 6329, 6329, ...
## Resampling results across tuning parameters:
##
##  k  Accuracy  Kappa
##  3  0.7512522  0.3332301
##  4  0.7495158  0.3265884
##  5  0.7636237  0.3495568
##  6  0.7629710  0.3468718
##  7  0.7714460  0.3628864
##  8  0.7711625  0.3617007
##  9  0.7791257  0.3774556
## 10  0.7770776  0.3691445
## 11  0.7814279  0.3769666
## 12  0.7803761  0.3714237
## 13  0.7825934  0.3727722
## 14  0.7821668  0.3697731
## 15  0.7828208  0.3706669
## 16  0.7827635  0.3693149
## 17  0.7826787  0.3668893
## 18  0.7831342  0.3660530
## 19  0.7846132  0.3698433
## 20  0.7841578  0.3680674
```

```

## 21 0.7843281 0.3677020
## 22 0.7843851 0.3662550
## 23 0.7843569 0.3646770
## 24 0.7847263 0.3653045
## 25 0.7852386 0.3651236
## 26 0.7852093 0.3639776
## 27 0.7851246 0.3633728
## 28 0.7840721 0.3593145
## 29 0.7839874 0.3580086
## 30 0.7832193 0.3549772
## 31 0.7832469 0.3545445
## 32 0.7827063 0.3518629
## 33 0.7821095 0.3502417
## 34 0.7819674 0.3483380
## 35 0.7827932 0.3499306
## 36 0.7817686 0.3458876
## 37 0.7824518 0.3473010
## 38 0.7827645 0.3474467
## 39 0.7827927 0.3466259
## 40 0.7830206 0.3464548
## 41 0.7831633 0.3465271
## 42 0.7825365 0.3427540
## 43 0.7827927 0.3438432
## 44 0.7823087 0.3414474
## 45 0.7814273 0.3376478
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 25.

```

```
topThreeModel
```

```

## k-Nearest Neighbors
##
## 7032 samples
## 3 predictor
## 2 classes: 'No', 'Yes'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 5 times)
## Summary of sample sizes: 6330, 6329, 6329, 6329, 6328, 6328, ...
## Resampling results across tuning parameters:
##
## k Accuracy Kappa
## 3 0.7810877 0.4101355
## 4 0.7818552 0.4122577
## 5 0.7816849 0.4112815
## 6 0.7826802 0.4132486
## 7 0.7835338 0.4140237
## 8 0.7843868 0.4151292
## 9 0.7846713 0.4154510
## 10 0.7849275 0.4161003
## 11 0.7849271 0.4159485
## 12 0.7851549 0.4166810
## 13 0.7849273 0.4160800
## 14 0.7854392 0.4169988

```




```

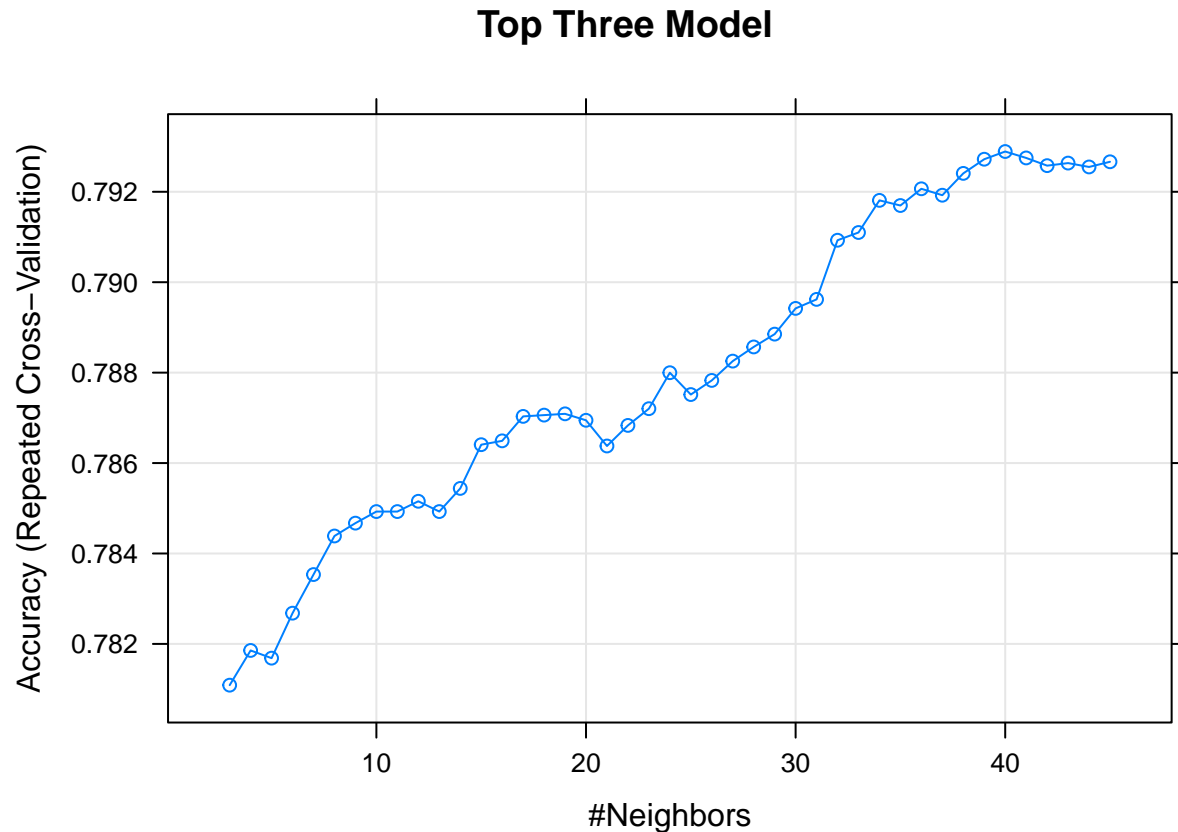
## 15 0.7864066 0.4188277
## 16 0.7864915 0.4182480
## 17 0.7870320 0.4185549
## 18 0.7870605 0.4168363
## 19 0.7870886 0.4149821
## 20 0.7869462 0.4131081
## 21 0.7863777 0.4105335
## 22 0.7868327 0.4101205
## 23 0.7872025 0.4092336
## 24 0.7879987 0.4100534
## 25 0.7875155 0.4075588
## 26 0.7878281 0.4080975
## 27 0.7882550 0.4088138
## 28 0.7885680 0.4094095
## 29 0.7888520 0.4095492
## 30 0.7894213 0.4106512
## 31 0.7896203 0.4099117
## 32 0.7909286 0.4115054
## 33 0.7910994 0.4107424
## 34 0.7918103 0.4119777
## 35 0.7916964 0.4108626
## 36 0.7920662 0.4118698
## 37 0.7919239 0.4119186
## 38 0.7924073 0.4137182
## 39 0.7927202 0.4150493
## 40 0.7928910 0.4155434
## 41 0.7927487 0.4150273
## 42 0.7925782 0.4144677
## 43 0.7926349 0.4145292
## 44 0.7925496 0.4141824
## 45 0.7926635 0.4145938
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 40.
plot(allVariableModel, main = "All Variable Model")

```

## All Variable Model



```
plot(topThreeModel, main = "Top Three Model")
```



### Model Outputs

We can see that each model uses a largely different value of K. Where the All Variable Model uses only 24 neighbours to achieve it's peak accuracy, model 2 uses 40. The plots demonstrate the increases and decreases in accuracy with the change in the number of neighbours.

## Model Comparison

The two models produce similar guideline accuracy results being within 2% of one another however the best strategy for testing their true accuracy is to create an AUC curve and compare the results.

```
# Create the prediction statistics that are used to calculate
# the AUC
model1Predictions <- getPredictions(allVariableModel, customerDataset,
  customerDataset$Churn)
model2Predictions <- getPredictions(topThreeModel, customerDataset,
  customerDataset$Churn)

# Create a dataframe of the prediction results. This df will
# allow us to access the classification results
model1Prediction_df <- createPrediction_df(allVariableModel,
  customerDataset, predictionType = "prob", oneClass = "Yes",
  zeroClass = "No")
model2Prediction_df <- createPrediction_df(topThreeModel, customerDataset,
  predictionType = "prob", oneClass = "Yes", zeroClass = "No")


# Calculate the AUC for each model
```

```

allVariableModelAUC <- as.numeric(getModelAUC(model = allVariableModel,
  dataset = customerDataset, outcomeColumn = customerDataset$Churn,
  type = "prob"))
topThreeModelAUC <- as.numeric(getModelAUC(model = topThreeModel,
  dataset = customerDataset, outcomeColumn = customerDataset$Churn,
  type = "prob"))

# Plot the ROC curves on the same graph for comparison
plotROCCurves(model1Prediction = model1Predictions, model2Prediction = model2Predictions,
  main = "Model Comparison", model1Name = sprintf("All Variable Model. AUC %.2f%%",
    allVariableModelAUC * 100), model2Name = sprintf("Top Three Model. AUC %.2f%%",
    topThreeModelAUC * 100))

```



```

# Create a confusion matrix for each model
model1ConfusionMatrix <- createConfusionMatrix(allVariableModel,
  customerDataset, customerDataset$Churn, oneClass = "Yes",
  zeroClass = "No")

model2ConfusionMatrix <- createConfusionMatrix(topThreeModel,
  customerDataset, customerDataset$Churn, oneClass = "Yes",
  zeroClass = "No")
model1ConfusionMatrix

##      outcomeColumn
##      No  Yes
## No  4826 1086

```

```
##    Yes  337  783
model2ConfusionMatrix

##      outcomeColumn
##      No  Yes
##    No  4736 1016
##    Yes  427  853

# Calculate the difference in the AUC
aucDifference <- (topThreeModelAUC - allVariableModelAUC) * 100
# Calculate the accuracy of accurately predicting yes
model1YesAccuracy <- (model1ConfusionMatrix[2, 2]/(model1ConfusionMatrix[2,
  1] + model1ConfusionMatrix[2, 2])) * 100
model2YesAccuracy <- (model2ConfusionMatrix[2, 2]/(model2ConfusionMatrix[2,
  1] + model2ConfusionMatrix[2, 2])) * 100
```

## Model Comparison Results

Looking at the model's AUC we can see that the top three model is still the most accurate overall, beating the all variable model by 0.53%. Where the top three model begins to falter is how accurately it predicts yes as it is only accurate 66.64% where the All Variable Model is correct when it predicts that a customer will churn 69.91%.